

K-OS Memory Management

Kai Chen

CS134c Spring 2003

Overview

- A data structure to keep track of available memory. (in pages)
 - Bitmap : one bit for each page
 - Stacks : pushes on free pages
- Virtual Memory with Paging
 - Page directory, page tables, page frames ...

Paging Basics

- **Page Boundary**
 - Block of memory 4Kb aligned. (starting at address where the lower 12 bits are 0)
- **Page Directory**
 - An array of 4-byte page table specifiers. (up to 4Kb)
- **Page Table**
 - An array of 4-byte page specifiers. One page table maps 4Mb memory and takes up to 4Kb to store.
- **Page Frame**
 - 4Kb of contiguous memory. Starting at page boundary.
- **Page directory and page table entries**
 - Where a page points to. And page attributes.

Page Directory/Table Entry

31 ... 12	11...9	8 ... 7	6	5	4 ... 3	2	1	0
address	Avail	Reserved	D	A	Reserved	U/S	R/W	P

- Address: (physical) 20 bits is enough. Last 12 bits always 0
- Avail: Can be used however we want
- D: dirty bit
- A: accessed bit
- U/S: user/superuser
- R/W: read/read and write
- P: present

Setting up Paging (example)

```
Unsigned long * page_dir = (unsigned long *) 0x9C000;  
unsigned long * page_table = (unsigned long *) 0x9D000;  
unsigned long addr = 0;
```

```
for(i=0; i<1024; i++){  
    page_table[i] = addr | 3;  
    addr += 4096;  
}
```

```
page_dir[0] = page_table | 3;
```

```
for(i=0; i<1024; i++){  
    page_dir[i] = 0 | 2;  
}
```

Enabling Paging


- Put the address of the page directory in CR3.
- Set bit 31 of CR0 to enable paging.

Virtual Memory

- **Physical Memory Organization**
 - Physical pages
 - Page usage count
- **Virtual Memory Organization**
 - Kernel Page Directory: 0x70000
 - Kernel Page Tables: 0x71000

Self Reference

- The last entry of the page directory points to itself.
- We can refer to page directory and page tables via virtual memory:
 - $\text{PAGE_SELF} = 1023$
 - $\text{PAGE_DIR_VADDR} = (\text{PAGE_SELF} \ll 22) \mid (\text{PAGE_SELF} \ll 12)$
 - $\text{PAGE_TABLE_VADDR} = (\text{PAGE_SELF} \ll 22)$



VM Components

- **Page Allocator:**
 - Find a free page in `page_use_count` vector
 - Return the physical address
 - Deallocation is the natural opposite
- **Memory Map (`mmap(vaddr, paddr, attr)`):**
 - Compute PDE and PTE for `vaddr`:
 - $\text{PDE}(\text{vaddr}) = \text{vaddr} \gg 22$
 - $\text{PTE}(\text{vaddr}) = \text{vaddr} \gg 12$
 - Page align `paddr`
 - Insert `paddr` into the page table


VM Components (Continued)

- **Low Level Allocation (void *morecore):**
 - Allocation done in pages (4Kb aligned)
 - Use `alloc_page()` to allocate physical pages
 - Use `mmap()` to map pages to virtual address
 - Return the start of the virtual address
 - Kernel Heap starts at `0x100000`
 - We do not allow down sizing

VM Components (Continued)

- **High Level Allocation (kmalloc(size_t size))**
 - Organize memory in blocks.
 - Each block has a header:

```
struct header {  
    size_t size;  
    header *next;  
    unsigned magic : 31  
    unsigned used : 1  
}
```



VM Components (Continued)

- **High Level Allocation: kmalloc (size_t size)**
 - Go through the kernel heap to find a large enough free block.
 - If found:
 - split it into 2 and save the extra
 - return the block with right size
 - If not found:
 - use morecore() to allocate a new block
 - add to kernel heap
 - split to save the extra
 - return the block with right size

VM Components (Continued)

- **High Level Deallocation: kfree()**
 - sanity check: consistent magic number
 - find the block in the heap
 - return error if not found
 - mark the block as free
 - combine the adjacent free blocks to reduce fragmentation

Summary

- **Allocation:**
 - allocate physical page
 - map to virtual memory
 - allocate virtual memory in 4Kb blocks
 - reorganize into linked list of finer blocks
- **Deallocation:**
 - mark block as free
 - combine adjacent free blocks